

# JUHUN PARK

202-924-4546 | [juhunpark32@gmail.com](mailto:juhunpark32@gmail.com) | [juhunpark.me](http://juhunpark.me) | [linkedin.com/in/juhun-park](https://linkedin.com/in/juhun-park) | [github.com/juhun32](https://github.com/juhun32)

## EDUCATION

### George Mason University

Bachelor of Science in Computer Science

Expected May 2027

GPA: 3.9

**Relevant Coursework:** Data Structures & Algorithms, Low-level Programming

**Honors:** \$20,000 Merit Scholarship Recipient

## EXPERIENCE

### Google Developer Groups

Lead Developer, George Mason University

Aug 2025 - Present

Fairfax, VA

- Architected a course scheduler with Go and PostgreSQL to ingest and manage 10,000+ university course sections.
- Integrated backtracking algorithm to solve the constraint satisfaction problem, generating conflict-free timetables.
- Engineered a data pipeline using Python and GitHub Actions to automatically scrape course catalog with 99.9% accuracy.
- Scaled with RabbitMQ and Redis for asynchronous processing, supporting 100+ concurrent users without degradation.

### Northwestern University

Jun 2025 - Aug 2025

Software Engineer Intern

Evanston, IL

- Built an alumni networking platform with SvelteKit and ASP.NET, serving 150 users with 99.9% uptime on Cloudflare.
- Engineered a PostgreSQL schema, optimized indexes, SQL functions on Supabase achieving p99 sub-5ms query latency.
- Implemented a multi-tier caching strategy using Cloudflare CDN and a Redis cache, offloading 95% of read requests.

## PROJECTS

### Copium.dev | [GitHub](#) | [copium.dev](http://copium.dev)

Go, TypeScript, SvelteKit, PostgreSQL, Algolia, GCP, AWS, Cloudflare

- Built an internship management platform with SvelteKit and PostgreSQL, processing 2000+ weekly applications.
- Delivered an 80% reduction in query latency for BigQuery analytics by leveraging CQRS architecture.
- Migrated media storage to Cloudflare R2, architecting a zero-egress pipeline without incurring AWS data transfer fees.
- Implemented fault-tolerant compensating transactions for consistency across 3 data stores with automated retries.

### Californian Temple | [GitHub](#) | [calple.date](http://calple.date)

Go, Next.js, Firestore, Docker, GCP Cloud Run

- Implemented an event pipeline using GCP Pub/Sub to ingest disparate streams into a unified Firestore state model.
- Developed a Go microservice to consume event streams, merging data real-time to predict and identify schedule conflicts.
- Achieved sub-100ms latency for event synchronization by optimizing Firestore listener patterns & gRPC communication.

### Sequantial | [GitHub](#)

Python, Kafka, gRPC, Docker, Git, AC/acsys (Assetto Corsa Internal API)

- Built a data pipeline to ingest real-time telemetry from Assetto Corsa, interfacing directly with the Shared Memory API.
- Architected a high-throughput telemetry pipeline using Kafka to decouple game loop ingestion from data processing.
- Leveraged Kafka's keyed partitioning to guarantee sequential ordering of events, ensuring accuracy despite network loss.

## TECHNICAL SKILLS

- Languages:** Python, TypeScript, JavaScript, Go, Java, SQL, C, C#
- Frameworks:** React, Node.js, Django, Flask, Next.js, SvelteKit, FastAPI, .NET, PostgreSQL
- Tools & Platforms:** Git, GCP, Vercel, Cloudflare, Pub/Sub, AWS, Linux, RabbitMQ, Docker, Redis, Kafka